

Databases

Lecture 12
Exam Problems

I Choose the correct answer(s) for the following multiple choice questions. Each question has at least one correct answer.

1. In a SELECT query:

- a. the SELECT clause can contain arithmetic expressions
- b. according to the conceptual evaluation strategy, ORDER BY is evaluated before GROUP BY
- c. HAVING can contain row-level qualification conditions
- d. DISTINCT eliminates duplicates from the answer set
- e. none of the above answers is correct.

2. The natural join operator $R_1 * R_2$ in the relational algebra:

- a. returns a relation whose schema contains only the attributes in R_1 that don't appear in R_2
- b. returns a relation whose schema contains all the attributes in R_1 and R_2 , with common attributes appearing only once
- c. returns 2 relation instances
- d. is not associative
- e. none of the above answers is correct.

3. In the ANSI-SPARC architecture (for a database system), a database can have:
- a. exactly one symbolic structure
 - b. several conceptual structures
 - c. several external structures
 - d. several physical structures
 - e. none of the above answers is correct.

4. Consider relation $S[\underline{A}, \underline{B}, \underline{C}, D, E, F, G, H]$ with:

- primary key $\{A, B, C\}$, no other candidate keys;
- functional dependencies that are known to hold over S : $\{F\} \rightarrow \{H\}$, $\{C\} \rightarrow \{E, G\}$;
- no repeating attributes.

a. S is not 1NF

b. S is 2NF

c. S is not BCNF

d. S is 3NF

e. none of the above answers is correct.

5. In a B-tree of order 8:

- a. a non-terminal node has at most 8 subtrees
- b. a non-terminal node with 7 values has 8 subtrees
- c. terminal nodes can be on different levels
- d. a non-terminal node with 7 values has 6 subtrees
- e. none of the above answers is correct.

6. Let α , β and γ be subsets of attributes in a relational schema. If $\alpha \rightarrow \beta$ and $\beta \rightarrow \gamma$, then by transitivity:

a. $\alpha \rightarrow \gamma$

b. $\gamma \rightarrow \alpha$

c. $\beta \rightarrow \alpha$

d. $\gamma \rightarrow \beta$

e. none of the above answers is correct.

7. Let RepairLog[RID, MechanicID, RollerCoasterID, RepairTime] be a table in a SQL Server database. RepairLog has 100.000 records and 2 indexes: a unique clustered index on RID and a non-clustered index on MechanicID without nonkey columns.

Consider the following query:

```
SELECT RID, MechanicID, RepairTime  
FROM RepairLog  
WHERE MechanicID = 7
```

If the execution plan contains an *Index Seek (NonClustered)*, it also contains a:

- a. *Clustered Index Scan*
- b. *Index Scan (NonClustered)*
- c. *Key Lookup (Clustered)*
- d. *Index Trick (NonClustered)*
- e. none of the above answers is correct.

8-10. Consider the relational schema $S[\underline{FK1}, \underline{FK2}, A, B, C, D, E]$, with primary key $\{FK1, FK2\}$. Answer questions 8-10 using the legal instance below:

FK1	FK2	A	B	C	D	E
1	1	a1	b1	c1	7	2
1	2	a_	b3	c1	5	2
1	3	a2	b1	c2	Null	2
2	1	a3	b3	c2	Null	100
2	2	a3	b3	c3	Null	100

8. Consider queries Q_1 and Q_2 :

Q_1 :

SELECT *

FROM S s1 LEFT JOIN S s2 ON s1.FK1 = s2.E

Q_2 :

SELECT DISTINCT *

FROM S s1 INNER JOIN S s2 ON s1.FK1 = s2.E

The cardinality of the answer set of Q_i is denoted by $|Q_i|$.

$|Q_1| - |Q_2|$ is:

a. 0

b. 3

c. 10

d. 2

e. none of the above answers is correct.

FK1	FK2	A	B	C	D	E
1	1	a1	b1	c1	7	2
1	2	a_	b3	c1	5	2
1	3	a2	b1	c2	Null	2
2	1	a3	b3	c2	Null	100
2	2	a3	b3	c3	Null	100

S[FK1, FK2, A, B, C, D, E]

9. Regarding the functional dependencies of S:

a. at least one of the following dependencies is not satisfied by the instance: $\{A\} \rightarrow \{B\}$, $\{FK1, FK2\} \rightarrow \{A, B\}$, $\{FK1\} \rightarrow \{A\}$

b. by examining the instance, we can conclude that at least one of the following

dependencies is specified on the schema S: $\{A\} \rightarrow \{B\}$, $\{FK1\} \rightarrow \{A, B\}$, $\{FK1\} \rightarrow \{A\}$

c. at least two of the following dependencies are not satisfied by the instance: $\{FK2\} \rightarrow \{A, B\}$, $\{A\} \rightarrow \{E\}$, $\{A, B\} \rightarrow \{E\}$, $\{FK1, FK2\} \rightarrow \{E\}$

d. by examining the instance, we can conclude that at least two of the following dependencies are specified on the schema S: $\{FK2\} \rightarrow \{A, B\}$, $\{A\} \rightarrow \{E\}$, $\{A, B\} \rightarrow \{E\}$, $\{B\} \rightarrow \{C, E\}$

e. none of the above answers is correct.

FK1	FK2	A	B	C	D	E
1	1	a1	b1	c1	7	2
1	2	a_	b3	c1	5	2
1	3	a2	b1	c2	Null	2
2	1	a3	b3	c2	Null	100
2	2	a3	b3	c3	Null	100

S[FK1, FK2, A, B, C, D, E]

10. Consider queries Q_1 and Q_2 :

Q_1 :

```
SELECT FK2, FK1, COUNT(DISTINCT B)
FROM S
GROUP BY FK2, FK1
HAVING FK1 = 0
```

Q_2 :

```
SELECT FK2, FK1, COUNT(C)
FROM S
GROUP BY FK2, FK1
HAVING MAX(E) < 0
```

The cardinality of the answer set of Q_i is denoted by $|Q_i|$.

FK1	FK2	A	B	C	D	E
1	1	a1	b1	c1	7	2
1	2	a_	b3	c1	5	2
1	3	a2	b1	c2	Null	2
2	1	a3	b3	c2	Null	100
2	2	a3	b3	c3	Null	100

$S[\underline{\text{FK1, FK2}}, \text{A, B, C, D, E}]$

->

10.

$|Q_1| - |Q_2|$ is:

a. 0

b. 2

c. 1

d. -1

e. none of the above answers is correct.

FK1	FK2	A	B	C	D	E
1	1	a1	b1	c1	7	2
1	2	a_	b3	c1	5	2
1	3	a2	b1	c2	Null	2
2	1	a3	b3	c2	Null	100
2	2	a3	b3	c3	Null	100

S[FK1, FK2, A, B, C, D, E]

11. A secondary index:

- a. can contain duplicates
- b. cannot contain duplicates
- c. can be non-clustered
- d. cannot be non-clustered
- e. none of the above answers is correct.

12. For the relation $R[A, B, C]$ below, consider the 3 possible projections on 2 attributes: $AB[A, B]$, $BC[B, C]$, $AC[A, C]$. How many extra records does $AB * BC * AC$ contain (i.e., records that don't appear in R)?

A	B	C
a1	b2	c1
a1	b1	c2
a2	b1	c1

- a. 0
- b. 1
- c. 2
- d. 3
- e. none of the above answers is correct.

13. The cross-product operator $R_1 \times R_2$ in the relational algebra:

- a. returns a relation whose schema contains all the attributes in R_1 followed by all the attributes in R_2
- b. returns a relation whose schema contains only the attributes in R_1
- c. returns 3 relation instances
- d. is associative
- e. none of the above answers is correct.

II Answer the following questions / solve the following problems.

1. Rewrite the CREATE TABLE statements below such that the following restriction is enforced: one T1 entity can be associated with any number of T2 entities, and one T2 entity can be associated with at most one T1 entity. Don't add other SQL statements.

```
CREATE TABLE T1  
(IDT1 INT PRIMARY KEY,  
C1 VARCHAR(100))
```

```
CREATE TABLE T2  
(IDT2 INT PRIMARY KEY,  
C2 DATE)
```

2. Write the relational algebra expression below as a SQL query.

$$\pi_{\{A,B,C\}}((\pi_{\{A,B,ID\}}(\sigma_{M=70}(R))) \otimes_{ID=IDT1} (\pi_{\{C,IDT1\}}(\sigma_{N>5}(S))))$$

3-6. Consider the relational schema $R[\underline{RID}, A, B, C, D, E, F]$, with primary key $\{RID\}$. Answer questions 3-6 using the legal instance below:

RID	A	B	C	D	E	F
1	100	200	5	200	20	11
2	101	50	11	200	5	12
3	100	100	7	200	5	13
4	200	200	6	200	20	14
5	200	100	2	200	5	9
6	300	50	11	200	5	10

3. What's the result set returned by the following query? Write the tuples' values and the names of the columns.

```
SELECT r1.RID, r1.A + r2.A C2, r1.A * r2.A C3
FROM R r1 LEFT JOIN R r2 ON r1.RID = r2.RID
WHERE r1.A > ANY (SELECT B
                  FROM R
                  WHERE C < 10)
```

RID	A	B	C	D	E	F
1	100	200	5	200	20	11
2	101	50	11	200	5	12
3	100	100	7	200	5	13
4	200	200	6	200	20	14
5	200	100	2	200	5	9
6	300	50	11	200	5	10

R[RID, A, B, C, D, E, F]

4. Evaluate the expressions below. π doesn't eliminate duplicates. What's the cardinality of T ?

$$S := \sigma_{F < 13}(R)$$

$$T := \pi_{\{S.RID, S.A\}}(S \otimes_{S.D=R.D} R)$$

RID	A	B	C	D	E	F
1	100	200	5	200	20	11
2	101	50	11	200	5	12
3	100	100	7	200	5	13
4	200	200	6	200	20	14
5	200	100	2	200	5	9
6	300	50	11	200	5	10

R[RID, A, B, C, D, E, F]

5. What's the result set returned by the following query? Write the tuples' values and the names of the columns.

```
SELECT R.*  
FROM  
  (SELECT r1.RID, r2.A, r3.B  
   FROM R r1 INNER JOIN R r2 ON r1.A = r2.B  
   INNER JOIN R r3 ON r2.B > r3.D  
   WHERE r1.F > 10) t  
RIGHT JOIN R ON t.RID = R.RID
```

RID	A	B	C	D	E	F
1	100	200	5	200	20	11
2	101	50	11	200	5	12
3	100	100	7	200	5	13
4	200	200	6	200	20	14
5	200	100	2	200	5	9
6	300	50	11	200	5	10

R[RID, A, B, C, D, E, F]

6. Write all functional dependencies F such that:

- F is satisfied by the current instance of R;
- the dependent of F is $\{D\}$;
- the determinant of F has a single column.

RID	A	B	C	D	E	F
1	100	200	5	200	20	11
2	101	50	11	200	5	12
3	100	100	7	200	5	13
4	200	200	6	200	20	14
5	200	100	2	200	5	9
6	300	50	11	200	5	10

R[RID, A, B, C, D, E, F]

7. Rewrite the expression below using only operators in the set $\{\sigma, \pi, \times, \cup, -\}$.

$$(R \bowtie_{R.ID=S.RID} S) \cap (T \bowtie_{T.ID=U.TID} U)$$

8. Let P, Q, R be 3 relations with schemas P[PID, P1, P2, P3], Q[QID, Q1, Q2, Q3, Q4, Q5], R[RID, R1, R2, R3], and E an expression in the relational algebra:

$$E = \pi_{\{P2, Q2, Q4, R3\}} (\sigma_{PID = Q1 \text{ AND } QID = R2 \text{ AND } P3 = 'Bilbo' \text{ AND } Q5 = 100 \text{ AND } R1 = 7} (P \times Q \times R))$$

Optimize E and draw the evaluation tree for the optimized version of the expression.